
NOTE: SOLUTION FOR KDD-CUP 2021 WIKIKG90M-LSC

TECHNICAL REPORT

Weiyue Su
Team PGL
Baidu Inc., Shenzhen
China
suweiyue@baidu.com

Zeyang Fang
Team PGL
Baidu Inc., Shenzhen
China
fangzeyang@baidu.com

Hui Zhong
Team PGL
Baidu Inc., Shenzhen
China
zhonghui03@baidu.com

Huijuan Wang
Team PGL
Baidu Inc., Shenzhen
China
wanghuijuan03@baidu.com

Siming Dai
Team PGL
Baidu Inc., Shenzhen
China
daisiming@baidu.com

Zhengjie Huang
Team PGL
Baidu Inc., Shenzhen
China
huangzhengjie@baidu.com

Yunsheng Shi
Team PGL
Baidu Inc., Shenzhen
China
shiyunsheng01@baidu.com

Shikun Feng
Team PGL
Baidu Inc., Shenzhen
China
fengshikun01@baidu.com

June 15, 2021

ABSTRACT

WikiKG90M in KDD Cup 2021 is a large encyclopedic knowledge graph, which could benefit various downstream applications such as question answering and recommender systems. Participants are invited to complete the knowledge graph by predicting missing triplets. Recent representation learning methods have achieved great success on standard datasets like FB15k-237. Thus, we train the advance algorithms in different domains to learn the triplets, including OTE, QuatE, RotatE and TransE. Significantly, we modified OTE into NOTE (short for Norm-OTE) for better performance. Besides, we use both the DeepWalk and the post-smoothing technique to capture the graph structure for supplementation. In addition to the representations, we also use various statistical probabilities among the head entities, the relations and the tail entities for the final prediction. Experimental results show that the ensemble of state-of-the-art representation learning methods could draw on each other's strength. And the manual feature engineering further improves the results.

Keywords Knowledge Graph Completion · Knowledge Embedding · Ensemble Method

1 Introduction

Knowledge graphs are directed multi-relational graphs about facts, usually expressed in the form of (h, r, t) triplets, where h and t represent head entity and tail entity respectively, and r is the relation between head entity and tail entity, e.g. *(Geoffrey Hinton, citizen of, Canada)*. Large encyclopedic knowledge graphs, like Wikidata [Vrandečić and Krötzsch, 2014] and Freebase [Bollacker et al., 2008], can provide rich structured information about entities and benefit a wide range of applications, such as recommender systems, question answering and information retrieval. However, large knowledge graphs usually face the challenge of incompleteness. For example, 71% of people in Freebase have no birth place and 75% have no nationality [Dong et al., 2014]. Therefore, predicting these missing facts in a knowledge graph is a crucial task, also named as knowledge graph completion task. We can see Figure 1 for a clear understanding.

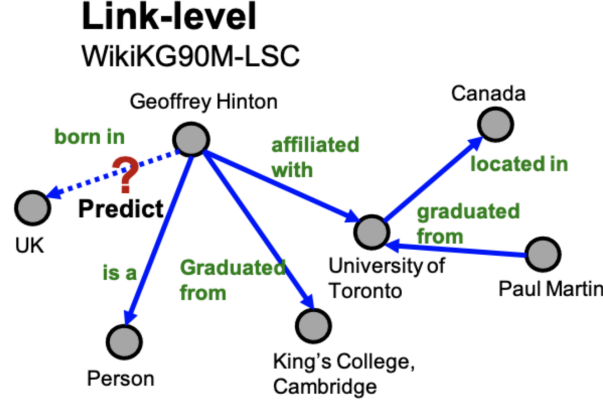


Figure 1: Impute missing triplets in knowledge graphs.¹

In order to address the issue of knowledge graph completion on large knowledge graphs, the 2021 KDD cup releases the WikiKG90M-LSC Task, which focuses on imputing missing facts in a knowledge graph extracted from the entire Wikidata knowledge base. Our method for this task consists of two stages: On the first stage, we propose an ensemble method for different knowledge embedding methods to build a strong model, in which, knowledge embedding methods, like TransE [Bordes et al., 2013] and RotatE [Sun et al., 2019], focus on embedding entities and relations into vectors and then we can use these embeddings to predict missing relations. On the second stage, we adopt several statistical features of the WikiKG90M dataset to help improve the final ensemble model performance. We conducted several experiments on WikiKG90M dataset to demonstrate the superiority of our method, and our team is currently among the awardees of the WikiKG90M-LSC track of the OGB-LSC, where we achieve 0.9727 MRR result in the final test set.

2 Methodology

2.1 Representation Learning

2.1.1 Triplet Embedding

The majority of knowledge graph representation algorithms relies on the triplets and they remain good interpretability in graph reasoning. In this competition, we adopt advance algorithms in different domains to encode the entities and relations, including NOTE, QuatE, RotatE and TransE. Considering the properties of each model, we ensemble their score results for the final prediction.

NOTE Tang et al. [2020] proposes OTE to model the symmetry/antisymmetry, inversion, and composition patterns. It takes the relations as an orthogonal transform in a high dimensional space. Relation matrix is orthogonalized so its inverse matrix can be obtained by simple transposing. The full model can be seen as an ensemble of K OTE models. The score functions are defined as Eq. 1 and 2.

$$d((h, r), t) = \sum_{i=1}^K (\|s_r^h(i)\phi(M_r(i))e_h(i) - e_t(i)\|), \quad (1)$$

$$d(h, (r, t)) = \sum_{i=1}^K (\|s_r^t(i)\phi(M_r(i))^T e_t(i) - e_h(i)\|), \quad (2)$$

where $s_r^h(i) = \text{diag}(\exp(s_r(i)))$ and $s_r^t(i) = \text{diag}(\exp(-s_r(i)))$ are the weights of relation matrix, ϕ is the Gram-Schmidt process.

However, though Tang et al. [2020] has scaled the L_2 norm of relation embeddings through scalar tensors $s_r(i) \in \mathcal{R}_{d_s}$, the convergence is still unstable in our experiments. The \exp operation in $s_r^h(i)$ and $s_r^t(i)$ enlarges values and could

¹<https://ogb.stanford.edu/assets/img/ogblsctaskoverview.png>

lead to this issue. Therefore, we further use L_2 norm to regularize the scalar tensors. In this case, the weights of relation matrix are modified into Eq. 3 and 4. We denote such modified version as NOTE (short for Norm-OTE).

$$s_r^h(i) = \frac{\text{diag}(\exp(s_r(i)))}{\|\text{diag}(\exp(s_r(i)))\|}, \quad (3)$$

$$s_r^t(i) = \frac{\text{diag}(\exp(-s_r(i)))}{\|\text{diag}(\exp(-s_r(i)))\|}. \quad (4)$$

QuatE Zhang et al. [2019] extend ComplEx [Trouillon et al., 2016] into quaternion space for better geometrical interpretations and more latent inter-dependencies. It proves that the Hamilton product in quaternion space can model both symmetry/antisymmetry and inversion patterns except the composition pattern.

RotatE RotatE can be view as an orthogonal transform in complex domain. Each relation is taken as a rotation from the head entity to the tail entity. Sun et al. [2019] has proved that the Hadamard product is able to model various relation patterns. In theory, it has equivalent model capability as OTE, so we also take it as a basic method.

TransE Bordes et al. [2013] interprets relationships as translations operating on entity embeddings in a real field. Although such assumption is not able to model complex relationships such as 1-N, N-1 and N-N, experiments on different datasets have proved its robustness and it can model composition pattern. Thus, we use results of this model to ensemble in our experiments.

2.1.2 Graph Context

Besides the specific triplets, the structure of sub-graphs comprised of triplets also reflect some semantic information. For example, *Geoffrey Hinton* in Figure 1 is connected with *King's College Cambridge* through relation *Graduated from*. We can infer that *Geoffrey Hinton* should be a person and is probable to be born in UK, as *King's College Cambridge* is located in UK. Entities and relations in this sub-graph influence each other. To supplement information of the graph structure, we use representative techniques such as Post-Smoothing [Klicpera et al., 2019] and DeepWalk [Perozzi et al., 2014].

Post-Smoothing We propose a two-stage method to capture the relationships among the entities in sub-graphs. Recent end-to-end models always use graph neural networks as encoder and triplet-based methods as decoder. But in our implementation, we take the opposite approach. In the first stage, we train TransE to encode the entities and relations according to the triplet context. In the second stage, we take the learned representations as entity embeddings. Then we propagate them through the entity adjacent matrix. A hyperparameter α decides the weight of the entity itself while $1 - \alpha$ denote weights of its neighborhood. The final updated embeddings are used for prediction. For example, for a given entity u , the final embedding u' is represented in Eq. 5.

$$u' = \alpha u + (1 - \alpha) \sum_{v \in \mathcal{N}(u)} v, \quad (5)$$

where u and v are the embeddings learned in the first stage.

DeepWalk Perozzi et al. [2014] propose DeepWalk to learn latent representations of nodes in homogeneous networks. In our solution, the relations in knowledge graph are ignored. We focus on the entity structure and use skip-gram technique to learn the semantic and structural correlations between entities along generated paths.

2.2 Manual Feature Engineering

In addition to embedding models, manual feature engineering is also a key part of our work. Since our goal is to predict tail through head and relation, our manual features include two parts, head to tail feature and relation to tail feature. Feature selection needs to be performed after obtaining the features.

2.2.1 Head to Tail Features

Head to tail feature is to predict the probability of tail by the current head. We start to walk from the head, and calculate the probability of walking through different paths to reach tail. For a head, relation and tail triple, it has 6 different walk




Figure 2: The walk paths of head to tail features. The yellow arrow is a direction of head entity to tail entity. HT is walk in the direction of head to tail, while TH is in direction of tail to head.

directions *direct*, including head to tail (HT), head to relation (HR), relation to head (RH), relation to tail (RT), tail to head (TH), tail to relation (TR). The probability of e_1 to e_2 in direction *direct* is as follow.

$$P_{direct}(e_1, e_2) = \frac{S_{direct}(e_1, e_2)}{\sum_{e \sim N_{direct}(e_1)} S_{direct}(e_1, e)} \quad (6)$$

$S_{direct}(e_1, e_2)$ is the frequency of e_1 to e_2 in the direction of *direct*. The frequency is calculated from all data, including the triplets and the tail candidates, respectively. We define 7 manual head to tail feature as follow.

$$F_{HT}(h, t) = P_{HT}(h, t) \quad (7)$$

$$F_{TH}(h, t) = P_{TH}(h, t) \quad (8)$$

$$F_{TH-HT}(h, t) = \sum_{e \sim N_{TH}(h) \cap N_{HT}(t)} P_{TH}(h, e) * P_{HT}(e, t) \quad (9)$$

$$F_{HT-HT}(h, t) = \sum_{e \sim N_{HT}(h) \cap N_{HT}(t)} P_{HT}(h, e) * P_{HT}(e, t) \quad (10)$$

$$F_{HT-TH}(h, t) = \sum_{e \sim N_{HT}(h) \cap N_{TH}(t)} P_{HT}(h, e) * P_{TH}(e, t) \quad (11)$$

$$F_{TH-TH}(h, t) = \sum_{e \sim N_{TH}(h) \cap N_{TH}(t)} P_{TH}(h, e) * P_{TH}(e, t) \quad (12)$$

$$F_{HT-HT-TH}(h, t) = \sum_{e_1 \sim N_{HT}(h)} \sum_{e_2 \sim N_{HT}(t)} P_{HT}(h, e_1) * P_{HT}(e_1, e_2) * P_{TH}(e_2, t) \quad (13)$$

2.2.2 Relation to Tail Features

We also define 5 manual relation to tail feature as follow.

$$F_{RT}(r, t) = P_{RT}(r, t) \quad (14)$$

$$F_{RH}(r, t) = P_{RH}(r, t) \quad (15)$$

$$F_{RT-TR-RT}(r, t) = \sum_{e_1 \sim N_{RT}(r)} \sum_{e_2 \sim N_{TR}(t)} P_{RT}(r, e_1) * P_{TR}(e_1, e_2) * P_{RT}(e_2, t) \quad (16)$$

$$F_{RH-HR-RT}(r, t) = \sum_{e_1 \sim N_{RH}(r)} \sum_{e_2 \sim N_{HR}(t)} P_{RH}(r, e_1) * P_{HR}(e_1, e_2) * P_{RT}(e_2, t) \quad (17)$$

$$F_{RT-HR-RT}(r, t) = \sum_{e_1 \sim N_{RT}(r)} \sum_{e_2 \sim N_{HR}(t)} P_{RT}(r, e_1) * P_{HR}(e_1, e_2) * P_{RT}(e_2, t) \quad (18)$$

2.2.3 Feature Selection

In order to combine the above-mentioned features and models, we use grid search for feature selection. The grid search will output the weights of the embedding models and manual features.

3 Experiments

3.1 Experimental Details

Original WikiKG90M dataset contains three time-stamps: September, October, and November of 2020, for training, validation, and testing, respectively, and only entities and relation types that appear in the earliest September knowledge graph are retained. The default parameter settings all models are batch_size=1000, learning rate of the mlp (mlp_lr)=2e-5, learning rate decay step (lrd_step)=1e-5, learning rate of the embedding (lr)=0.1, gamma=12 and hidden_size=200. Specifically, ote_size=20 in NOTE model. For our final submission, we mix the training data and validation data.

3.2 Experimental Results

Table 1 shows the specific structure of all the models we use for ensemble, and report the final validation and test results of our ensemble method.

Table 1: Experimental Results: "-" means default parameter setting.

Model Settings	Main Parameter Settings	Origin Validation Acc.
NOTE-1	-	0.9189
NOTE-2	batch_size=1200	0.9191
NOTE-3	lrd_step=2e5	0.9172
NOTE-4	lrd_step=4e4	0.9170
NOTE-5	mlplr=4e-5	0.9179
NOTE-6	lr=0.3	0.9181
NOTE-7	gamma=10	0.9160
NOTE-8	gamma=14	0.9173
NOTE-9	ote_size=40	0.9201
NOTE-10	hidden_size=240	0.9182
TransE	-	0.8559
TransE	post-smoothing	0.8800
RotatE	-	0.8871
RotatE	post-smoothing	0.8940
QuatE	-	0.8846
DeepWalk	-	0.6132
Ensemble the Above Models	Without manual features	0.9435
Ensemble the Above Models	With manual features	0.9797

4 Conclusion

In this paper, we present our solution for the final test. First, we ensemble different models for representation learning. Specifically, we propose the NOTE model to make the training process steady. And the DeepWalk and the post-smoothing technique are used to capture the graph structure information among learned embeddings. We also use recent advance models including QuatE, RotatE, TransE. Second, we detail the manual feature engineering. The selected features are used to adjust the predictions. The experimental results show that our solution achieves excellent performance on the WikiKG90M dataset.

References

- Denny Vrandečić and Markus Krötzsch. Wikidata: a free collaborative knowledgebase. *Commun. ACM*, 57(10):78–85, 2014.
- Kurt D. Bollacker, Colin Evans, Praveen Paritosh, Tim Sturge, and Jamie Taylor. Freebase: a collaboratively created graph database for structuring human knowledge. In *Proceedings of the ACM SIGMOD International Conference on Management of Data, SIGMOD 2008, Vancouver, BC, Canada, June 10-12, 2008*, pages 1247–1250. ACM, 2008.
- Xin Dong, Evgeniy Gabrilovich, Jeremy Heitz, Wilko Horn, Ni Lao, Kevin Murphy, Thomas Strohmman, Shaohua Sun, and Wei Zhang. Knowledge vault: A web-scale approach to probabilistic knowledge fusion. In *Proceedings of the 20th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 601–610, 2014.
- Antoine Bordes, Nicolas Usunier, Alberto García-Durán, Jason Weston, and Oksana Yakhnenko. Translating embeddings for modeling multi-relational data. In *Advances in Neural Information Processing Systems 26: 27th Annual Conference on Neural Information Processing Systems 2013. Proceedings of a meeting held December 5-8, 2013, Lake Tahoe, Nevada, United States*, pages 2787–2795, 2013.
- Zhiqing Sun, Zhi-Hong Deng, Jian-Yun Nie, and Jian Tang. Rotate: Knowledge graph embedding by relational rotation in complex space. In *7th International Conference on Learning Representations, ICLR 2019, New Orleans, LA, USA, May 6-9, 2019*, 2019.
- Yun Tang, Jing Huang, Guangtao Wang, Xiaodong He, and Bowen Zhou. Orthogonal relation transforms with graph context modeling for knowledge graph embedding. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, ACL 2020, Online, July 5-10, 2020*, pages 2713–2722, 2020.
- Shuai Zhang, Yi Tay, Lina Yao, and Qi Liu. Quaternion knowledge graph embeddings. In *Advances in Neural Information Processing Systems 32: Annual Conference on Neural Information Processing Systems 2019, NeurIPS 2019, December 8-14, 2019, Vancouver, BC, Canada*, pages 2731–2741, 2019.
- Théo Trouillon, Johannes Welbl, Sebastian Riedel, Éric Gaussier, and Guillaume Bouchard. Complex embeddings for simple link prediction. In *Proceedings of the 33rd International Conference on Machine Learning, ICML 2016, New York City, NY, USA, June 19-24, 2016*, volume 48, pages 2071–2080, 2016.
- Johannes Klicpera, Aleksandar Bojchevski, and Stephan Günnemann. Predict then propagate: Graph neural networks meet personalized pagerank. In *7th International Conference on Learning Representations, ICLR 2019, New Orleans, LA, USA, May 6-9, 2019*, 2019.
- Bryan Perozzi, Rami Al-Rfou, and Steven Skiena. Deepwalk: online learning of social representations. In *The 20th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, KDD '14, New York, NY, USA - August 24 - 27, 2014*, pages 701–710, 2014.